

 ARCHITECTURE & SCALING SUB-WORKING GROUP

# YOU ONLY TRAIN ONCE: MAKING ARCHITECTURAL DECISIONS FOR A >100B MODEL

Big Science Episode #2 — INLG, 2021/09/20



Big Science

Iz Beltagy

AI2

Julien Launay

Lighton

# The Scaling and Architecture Sub-Working Group

**What?**

Draft and validate an **architecture & training setup**  
to get the best out of our GPU budget.

**How?**

By establishing **principled baselines**,  
carefully **evaluating novel modelling choices**,  
and studying the **scaling of candidate architectures**.

**Constraints.**



**proven**

no unnecessary risks



**scalable**

final run: >200B param., 4MGPUh



**efficient**



**multilingual**



**emergent**

few-shot, prompt tuning, etc.

# Main unknowns in 🌸 Big Science



## Scale

**Very few models have been trained in the 100-200B range.**

GPT-3 (English, OpenAI), Jurassic-1 (English, A21),

HyperClova (Korean, Naver), PanGu-Alpha (Chinese, Huawei).

😊 with engineering working group.



## Multilinguality

**Limited knowledge on extreme-scale generative multilingual models.**

Closest comparison: mT5, 100 languages, 11B parameters. No large generative-only model.

**Can we avoid the *curse of multilinguality*?**

Severely underperforming monolingual counterparts.

😊 with multilingual working group.



## Architecture

**Bridge the LM and encoder-decoder performance gap with prefix LM.**

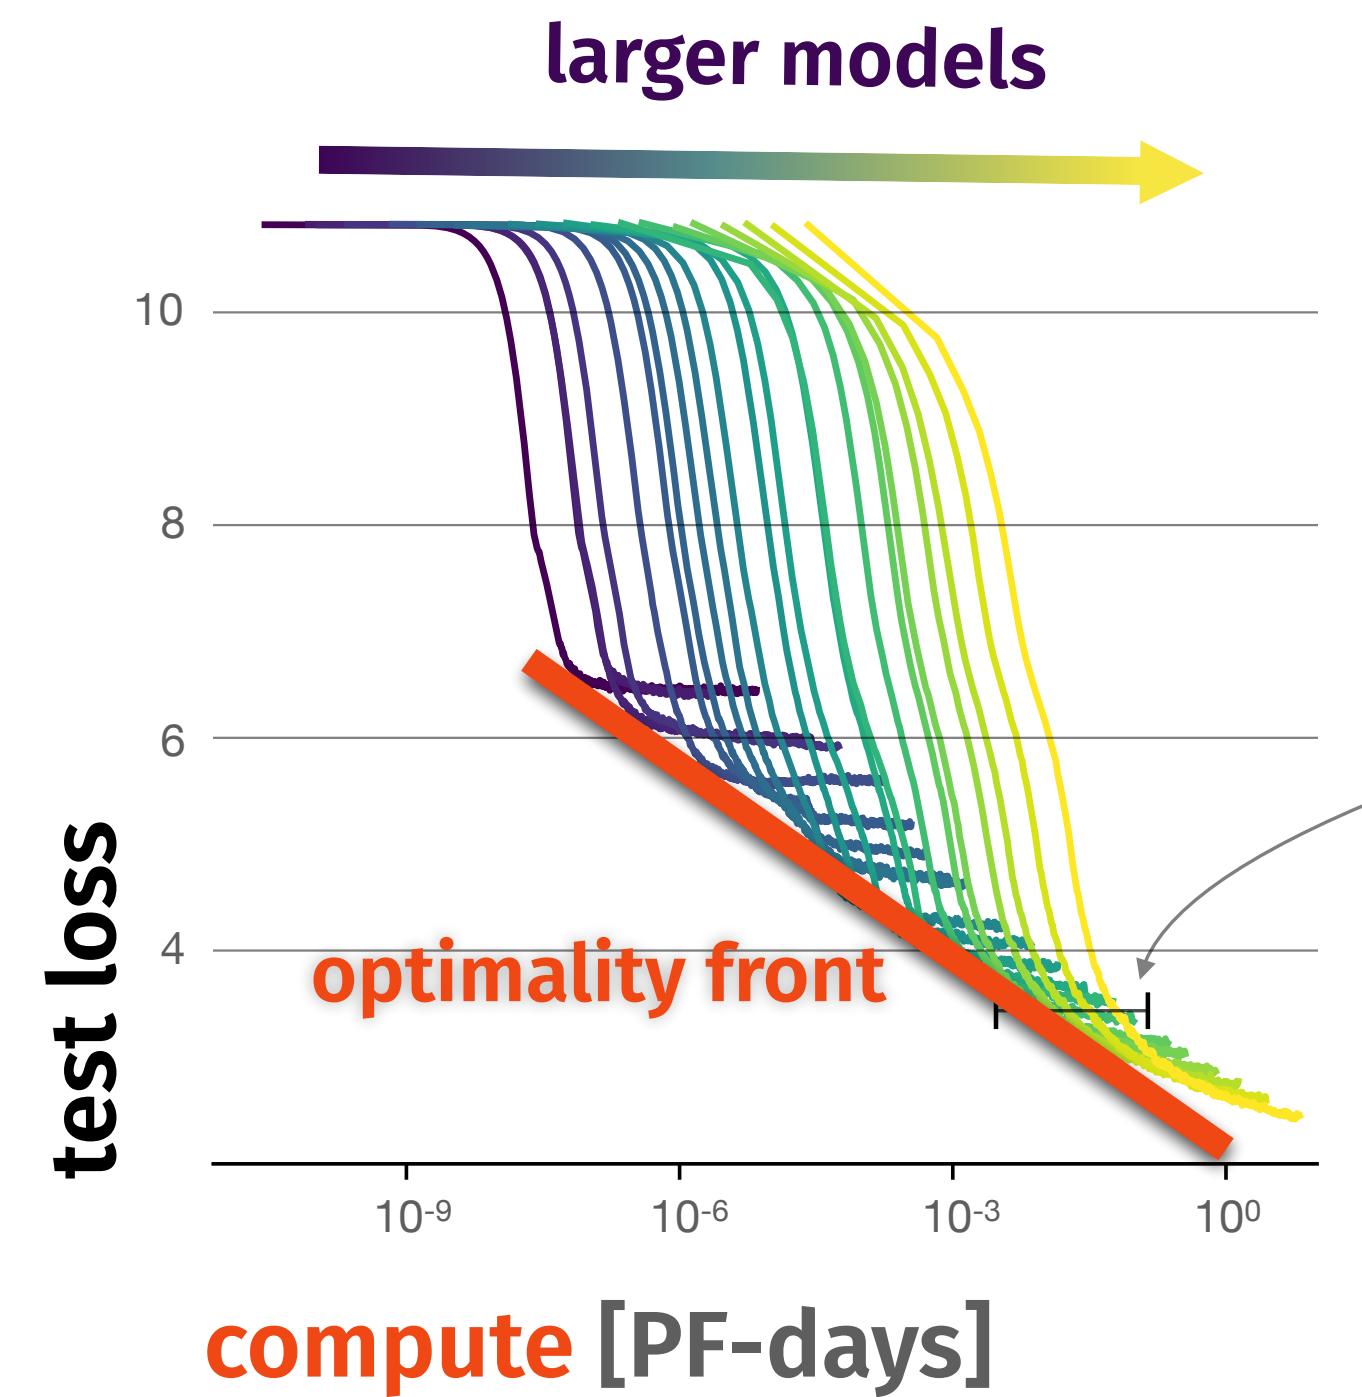
How to validate prefix LM at scale?

# Evaluations and metrics to benchmark architectures

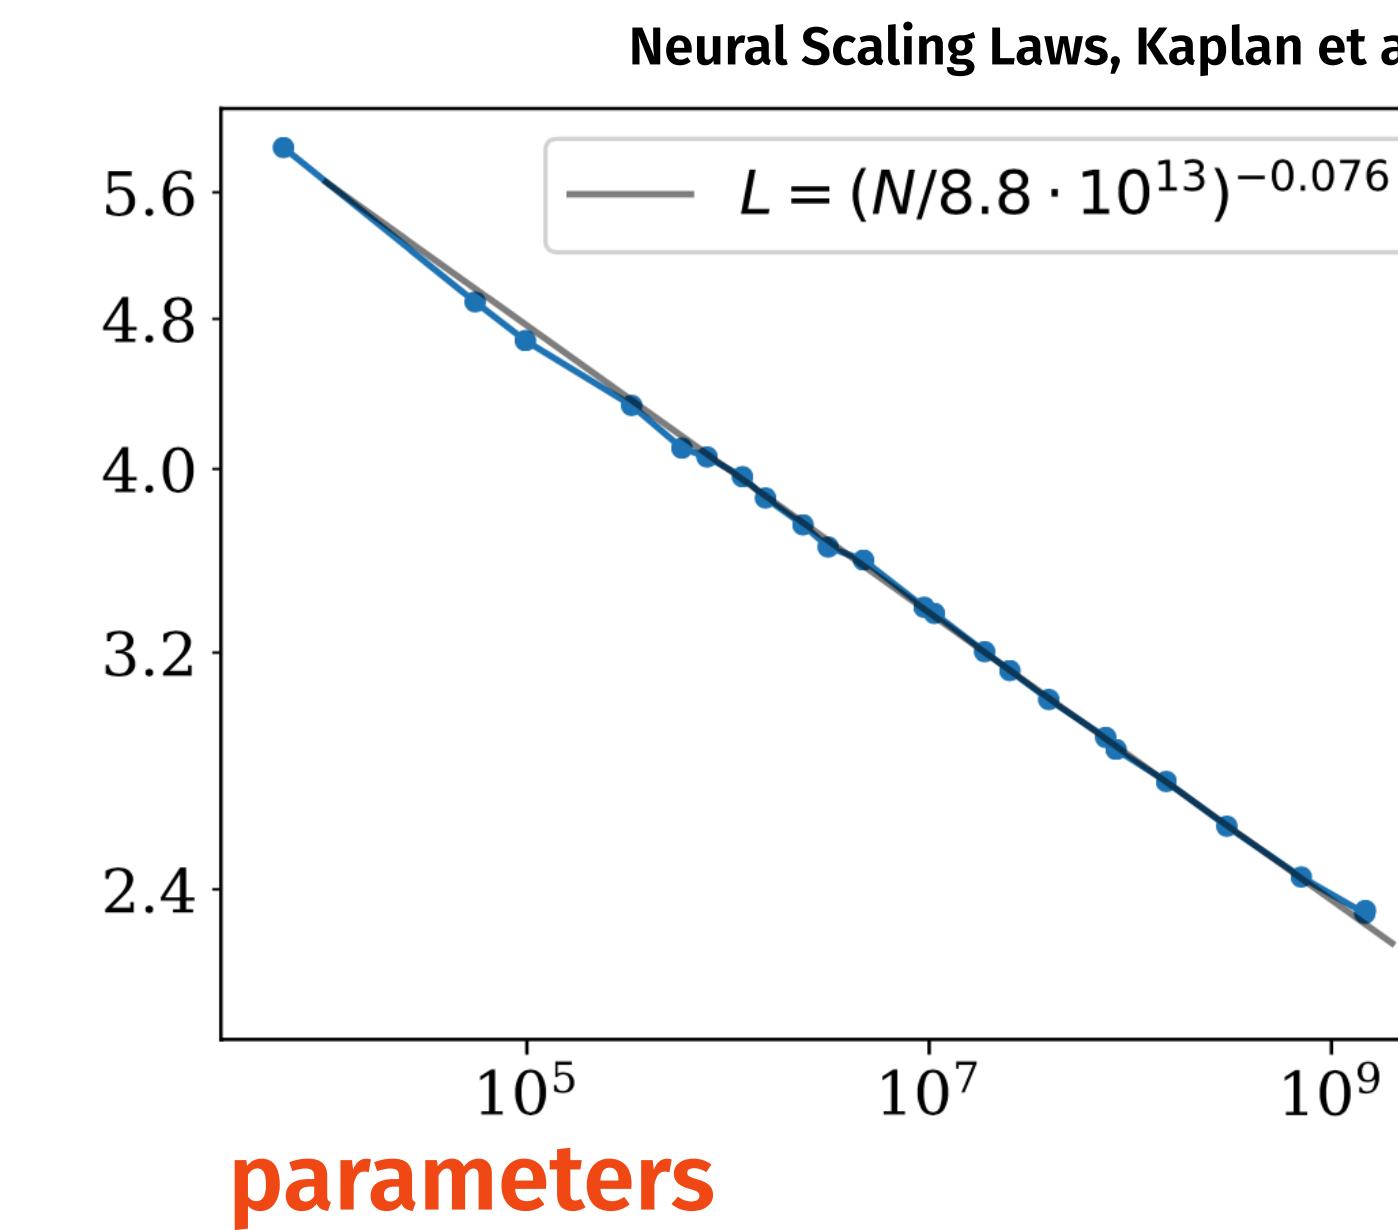
• Usual and simple metrics: validation loss, training time/throughput, etc.

efficiency & stability are key metrics at the 100B+ scale!

↗ Empirical backing of scaling laws to evaluate scaling:



performance is predictable  
using simple power laws



still, some behaviours are scale-emergent → train as large as possible, ~1B scale at least

## Evaluations and metrics to benchmark architectures



**Zero/few-shot performance evaluation on a large range of datasets.**

currently using Eleuther AI evaluation harness for English baselines.

😊 with evaluation group → multilingual evaluation, etc.



**Big unknown: how will final 200B model be used by the community?**



**Weights offloading/streaming make inference “accessible”...**

---

ZeRO-infinity

**but still very expensive to run in practice!** 💰



**Currently, OpenAI/A21/Cohere → hosted API with a text/log-prob interface.**

fine-tuning only offered for small models.



**Other approaches: efficient fine-tuning, adapter, prompt tuning, etc.**

keep emergent possibilities open!

## Unknown #1: Scale

📈 100B+ scale is **unforgiving**: we need excellent tooling, scalable architecture, etc.

every FLOP counts!

💥 “**Unstable**” behaviour in training at scale, not fully explained.

numerical instabilities: float16, etc. → can be avoided with bfloat16 on modern hardware (TPUs/A100s)

data-related instabilities? → see work on curriculum learning

diagnostic tools? → gradient noise scale, weightwatcher, etc.



🤖 Engineering working group: “big” **exploratory runs** at the >10B scale.

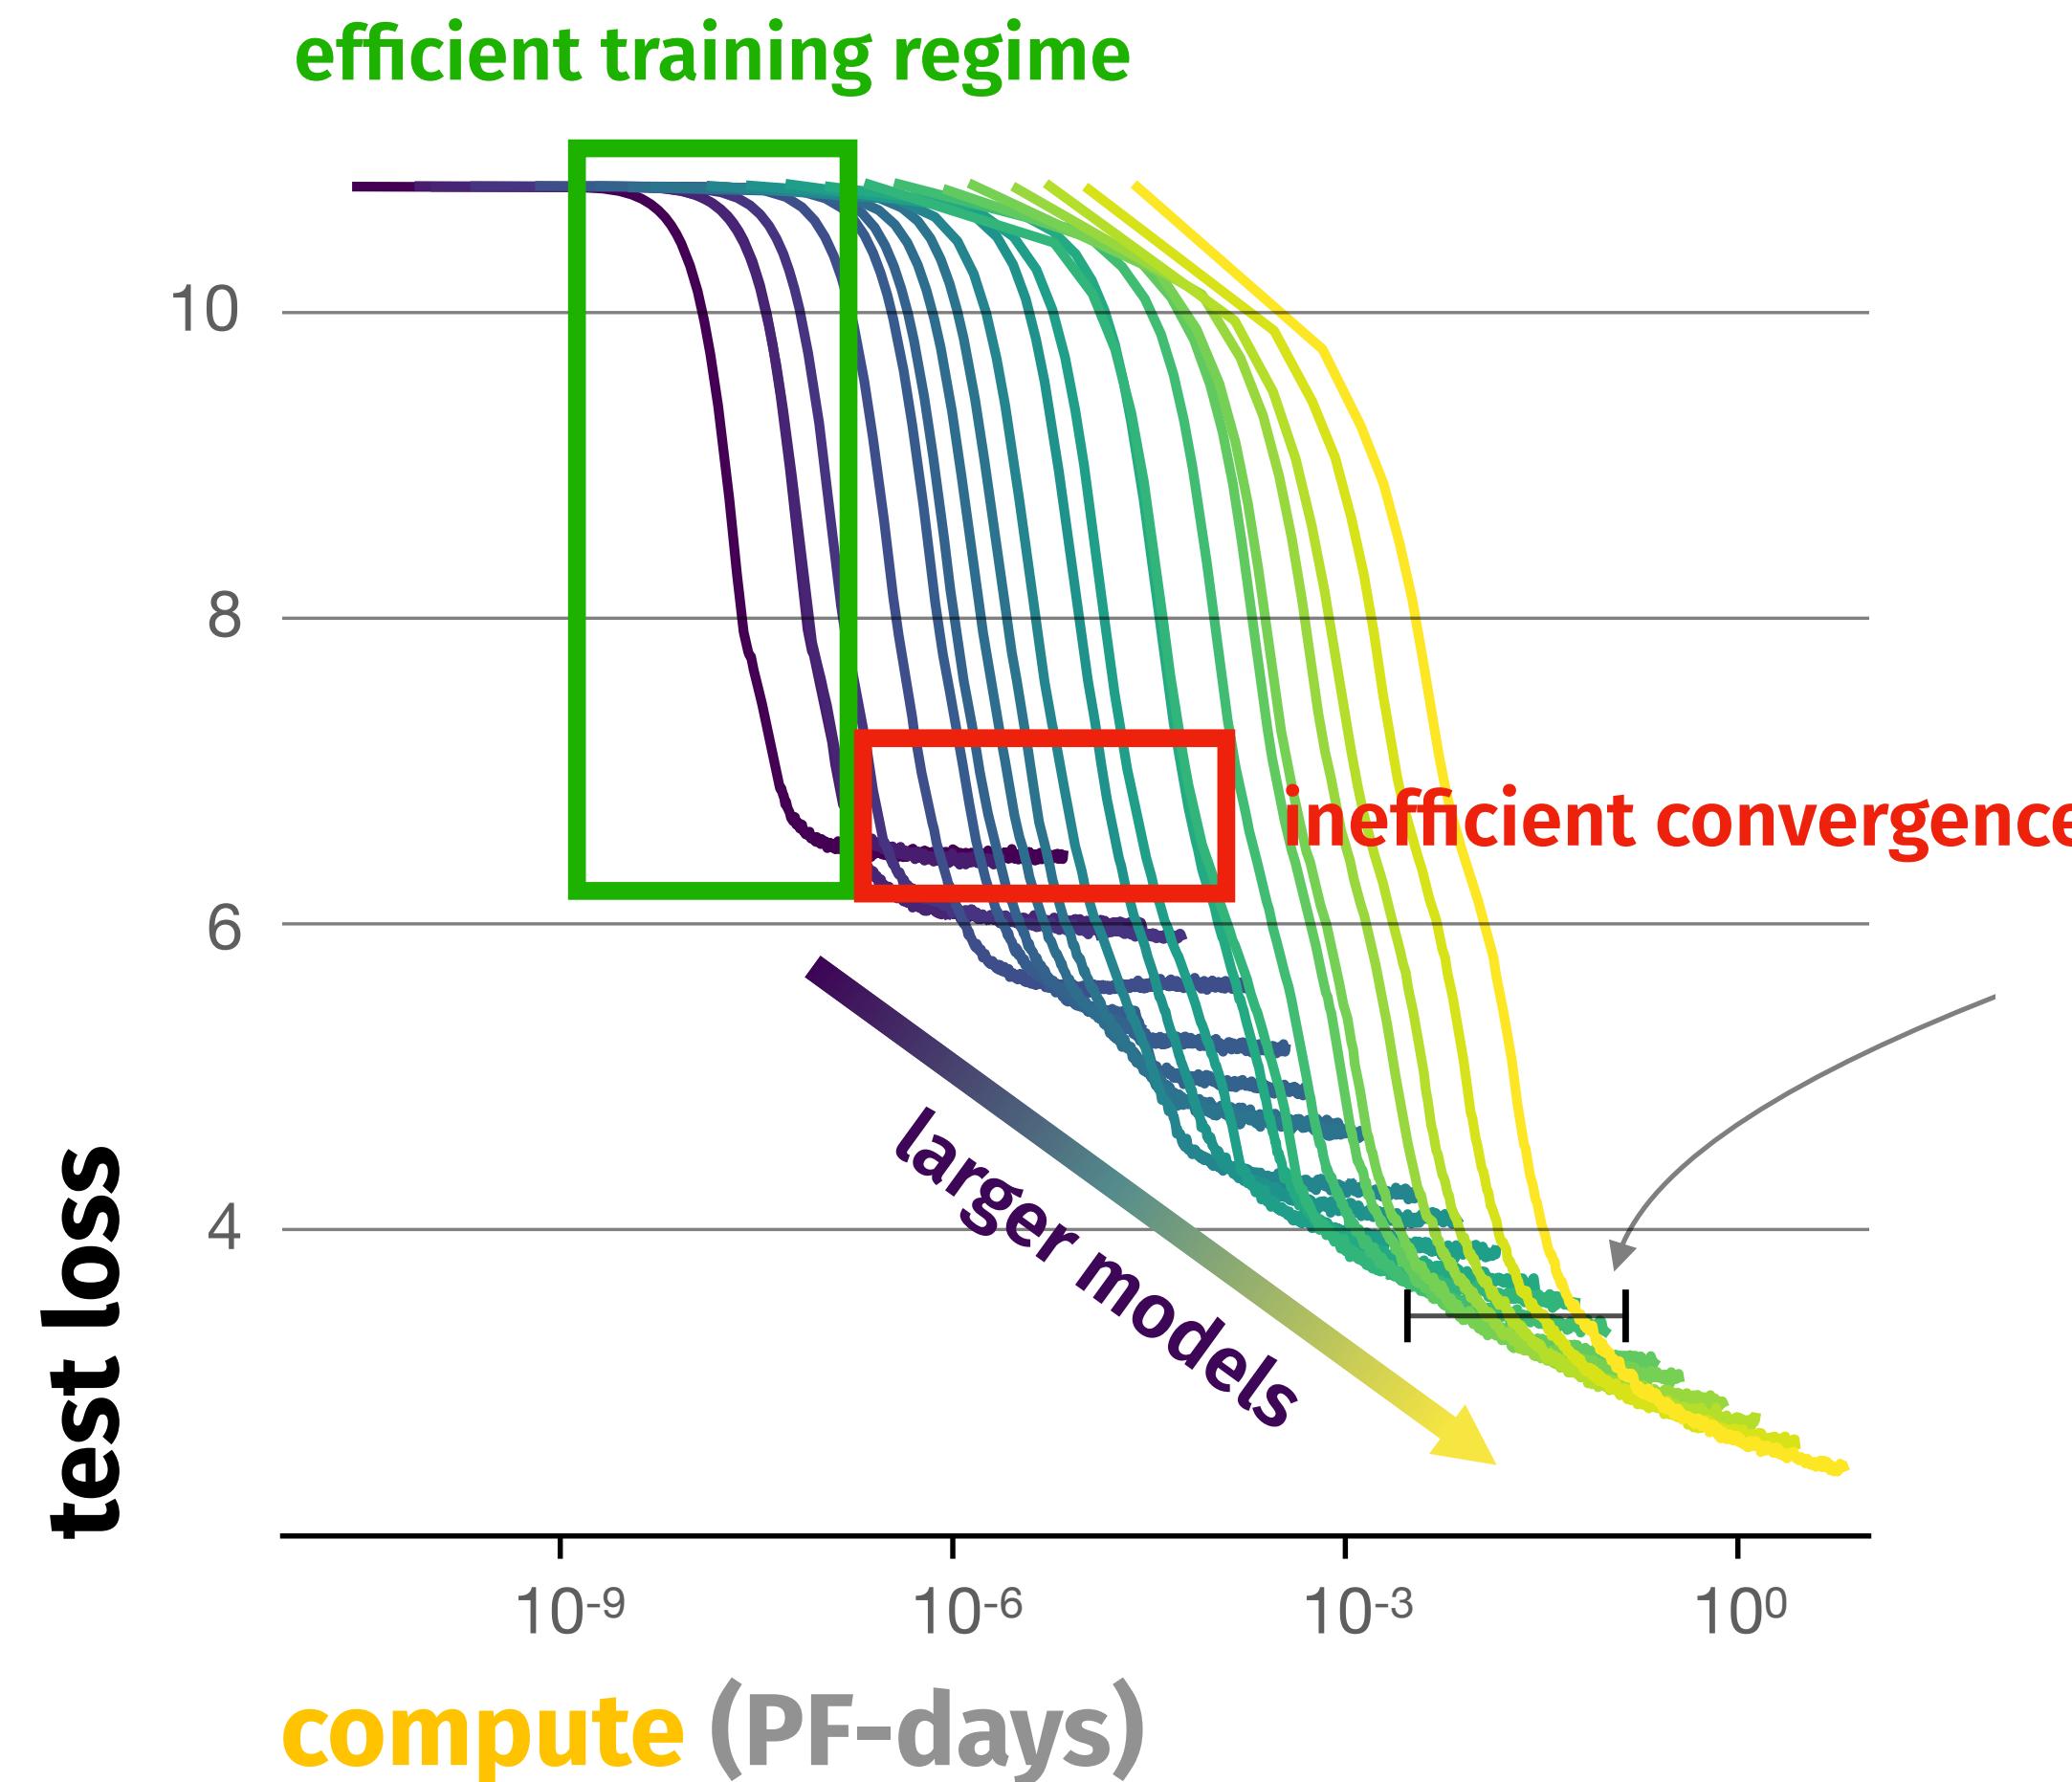
training #1 (13B English-only) complete, now looking at 13B multilingual for training #2.

one lesson already: dataset matters *a lot* for end-task performance!

# Training setup: at scale, training to convergence vs optimality

🧠 Don't train to convergence, but to optimality for efficiency in final run.

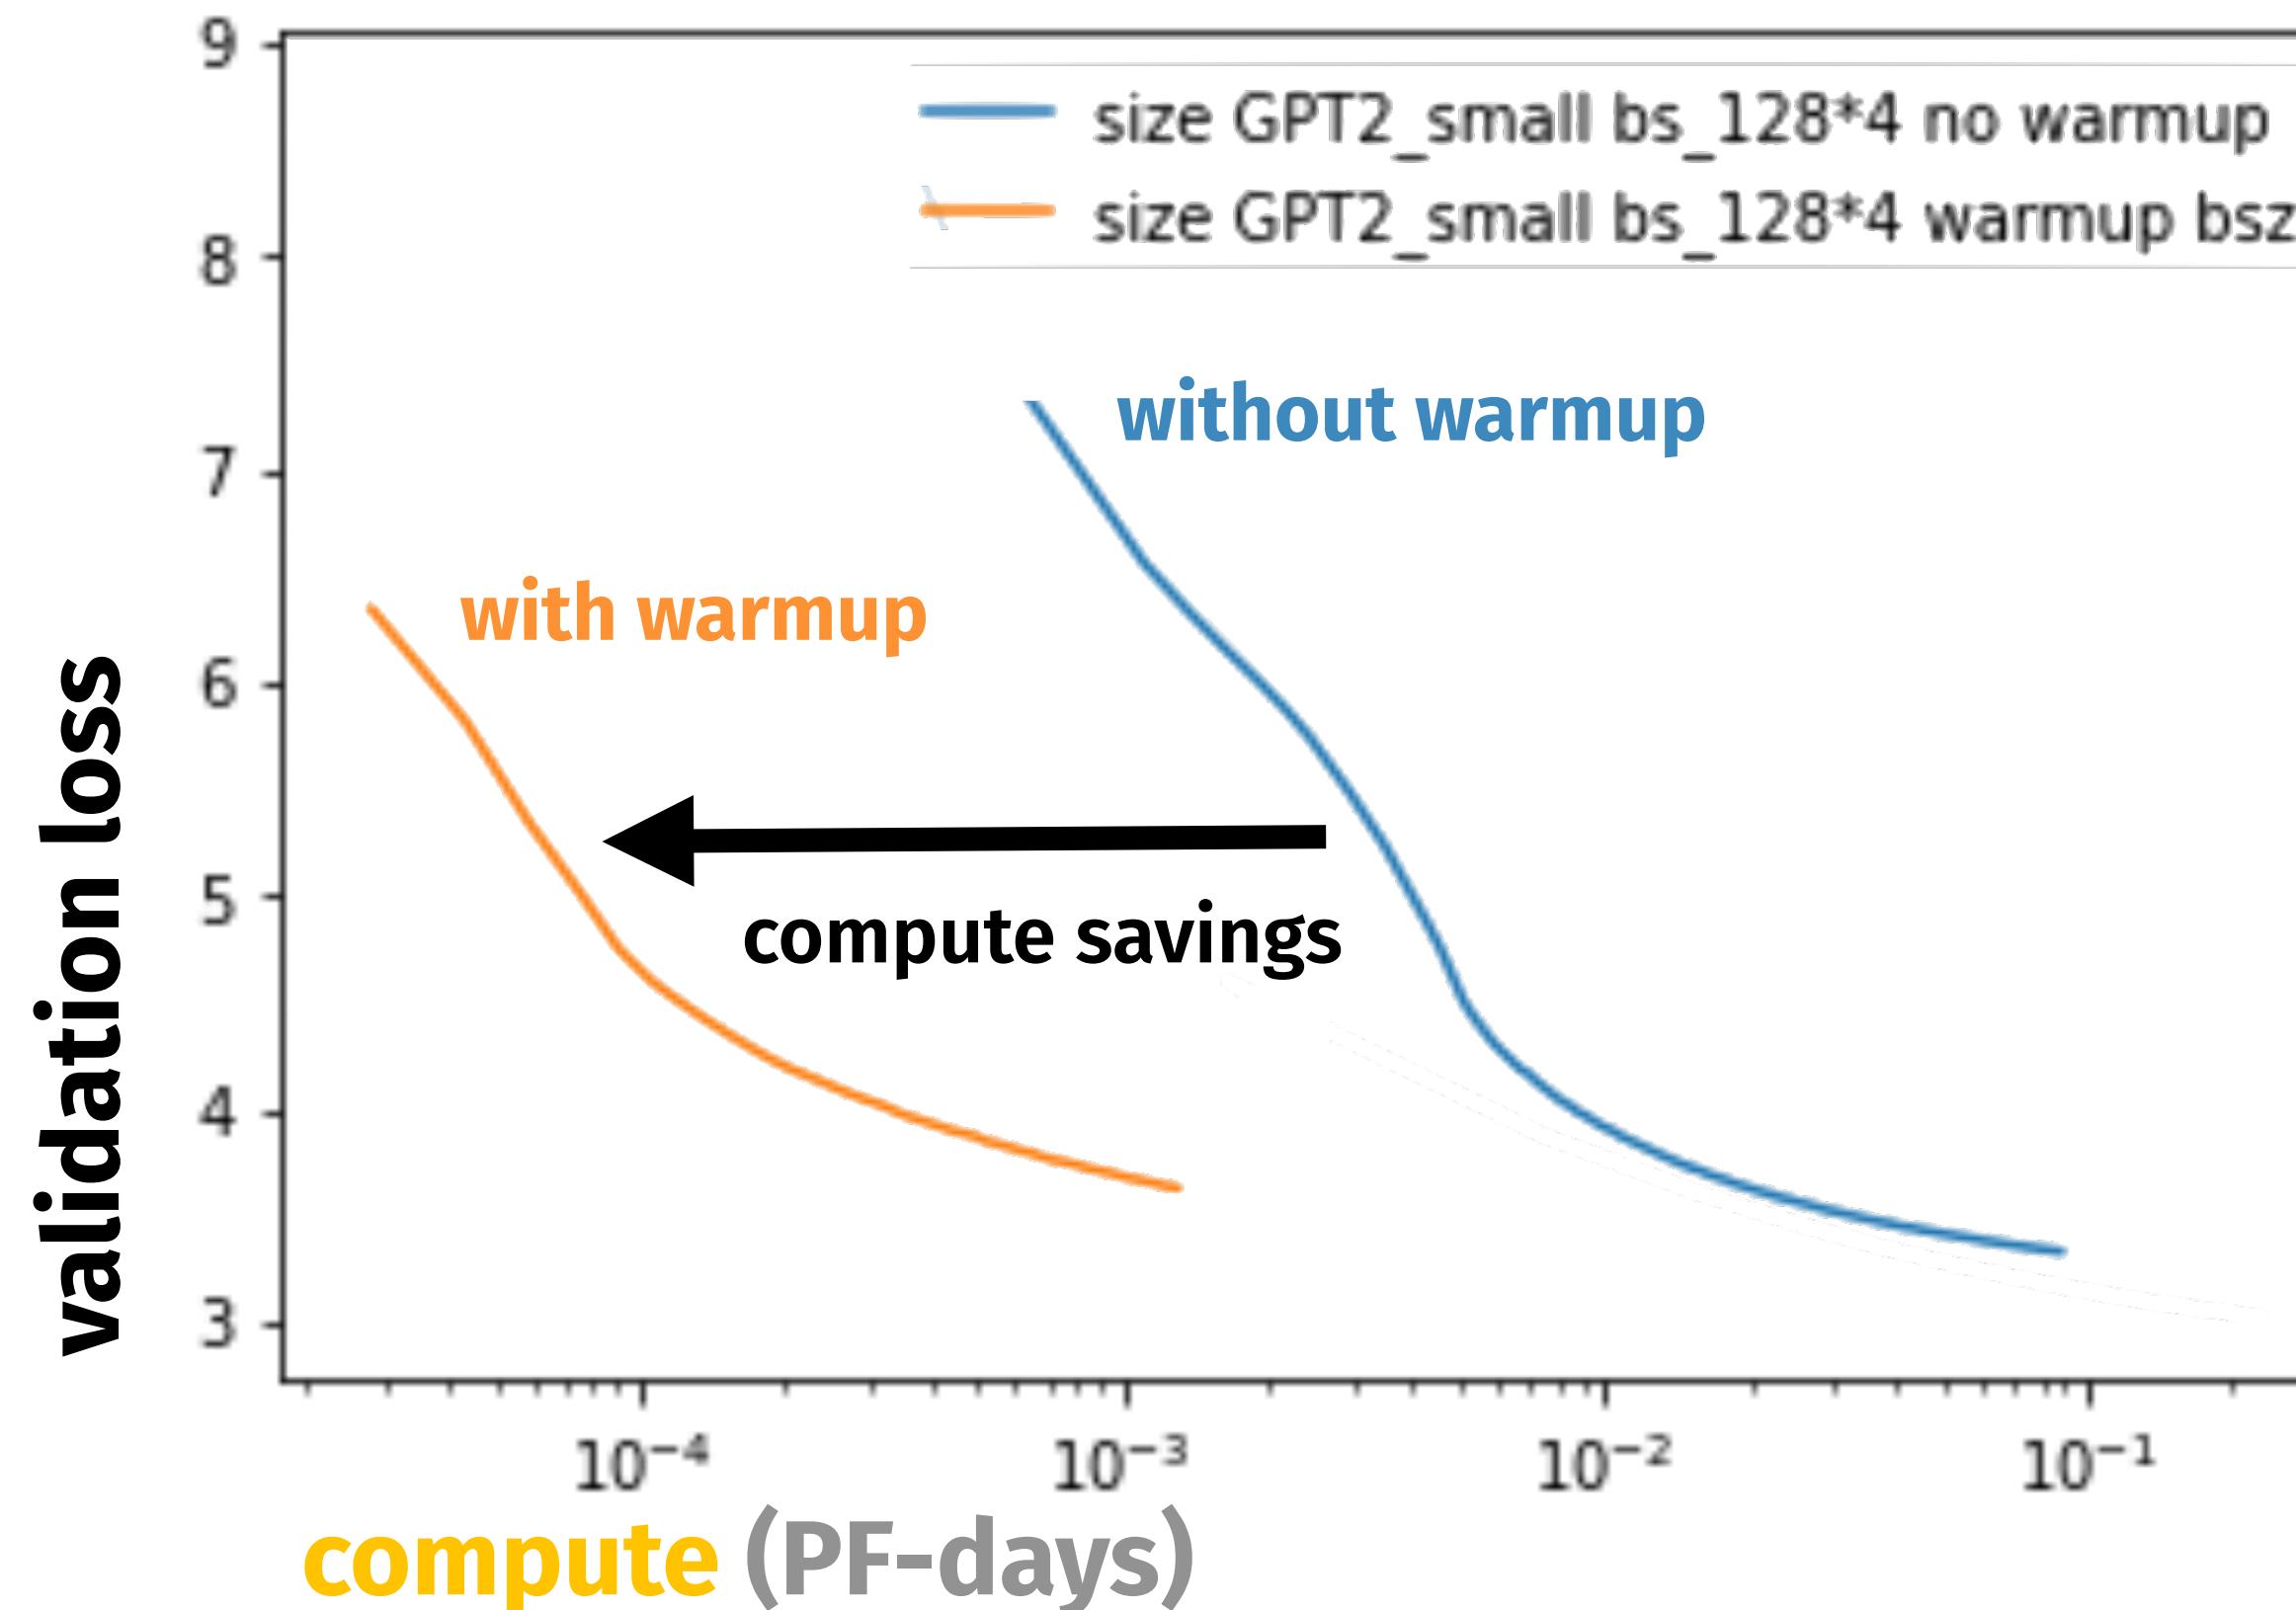
training budget: 200B, 4,400 PF-days (~4 MV100h@25 TFLOPs) to optimality, 30,000 PF-days (~30 MV100h) to conv.



# Batch size warmup saves compute

💡 **Batch size warmup:** start with a small batch size, then linearly increase to max batch size.

🧐 Intuition: **gradient noise** is high early in training, so large batch size is wasteful.



# Scaling laws as a **diagnostic** tool

🤔 Big Science training #1 (13B, English-only): disappointing few-shot performance.

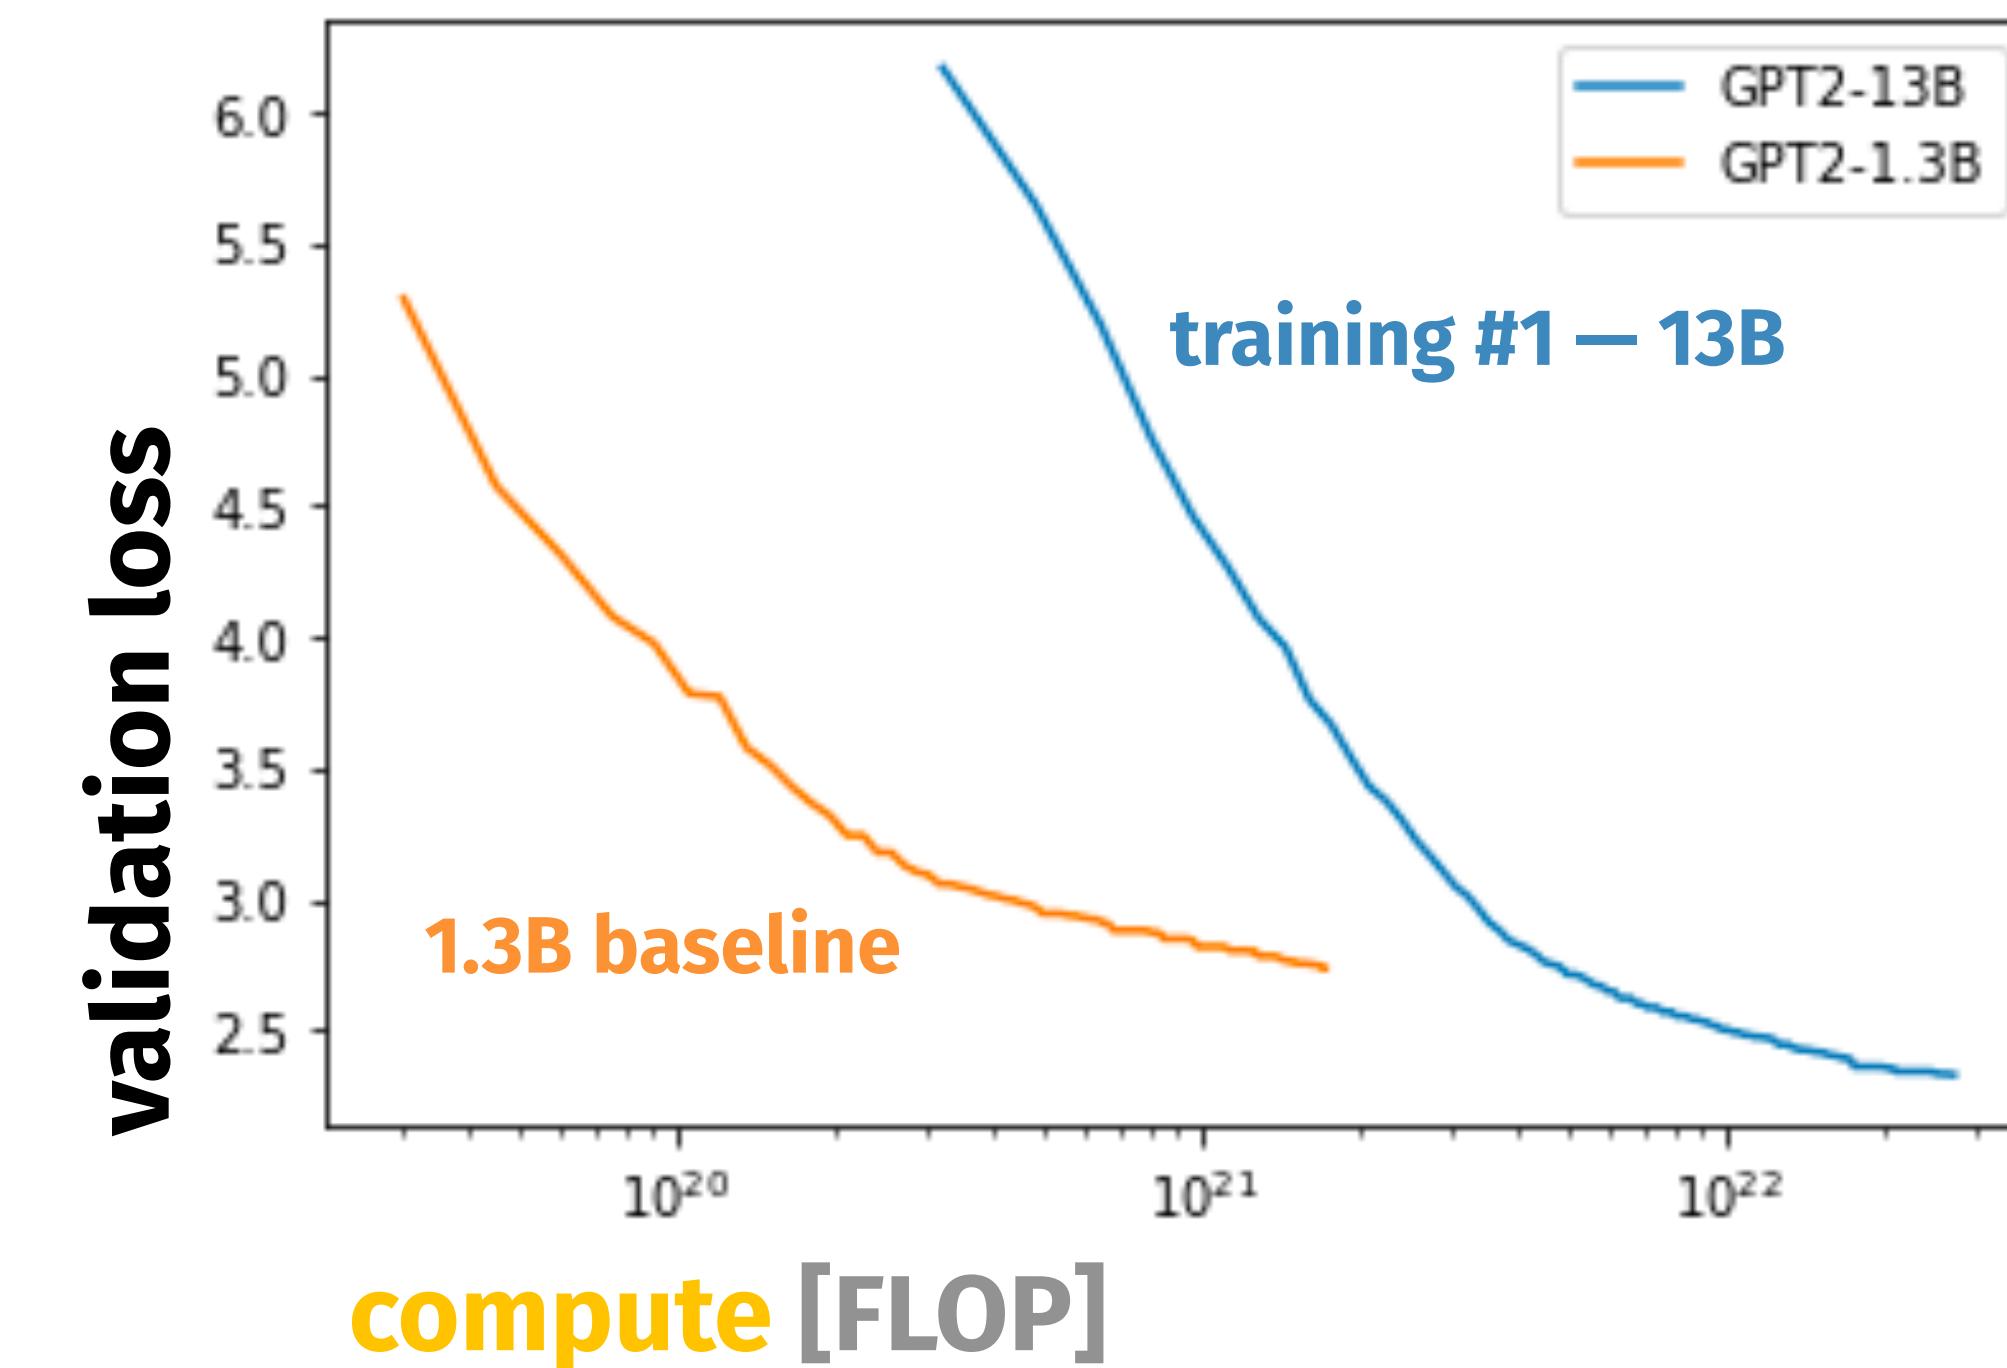
Results from EAI harness obtained by Stella Biderman.

<b>5.02</b>	<b>65%</b>	<b>54%</b>
lambada ppl.	winogrande acc.	hellaswag acc.

more in line with a **2.7-6B model!**

is this a data (**OSCAR**) problem?

or a setup problem?



## Unknown #2: Multilinguality

😊 Build a model that is **valuable** to the community at large.

languages selection, data collection, release licenses, etc. ————— many other WGs in Big Science!

🤔 Under-explored at scale, with ***curse of multilinguality*** problem.

if multilingual model severely underperforms monolingual counterparts, not that interesting!

no large-scale generative multilingual model exists... ————— ! very sensitive to data,  
no high-quality multilingual dataset!

100B English tokens vs 100B multilingual tokens, what's the gap?

😓 Evaluation of multilingual models is more challenging.

less big and “wide” benchmarks than in English for low-ressources languages.

# Tackling multilinguality under the angle of **scaling laws**



## Can we establish **multilingual scaling laws?**

quantify how languages scale differently...

quantify benefits from one language to another, like has been done for multimodal setups...

connect to fundamental linguistics works and validate findings



## Can we use this law for more **principled multilingual training**.

inform sampling strategy/scaling of gradients, etc.

We will be answering this questions soon 😊

## Unknown #3: Architecture

**GPT-3** as our base architecture, however...

⚠ From the T5 paper: performance of autoregressive LM is lower than encoder-decoder

Architecture	Objective	Params	Cost	GLUE	CNNDM	SQuAD	SGLUE	EnDe	EnFr	EnRo	T5, Raffel et al.
Encoder-decoder	Denoising	$2P$	$M$	<b>83.28</b>	<b>19.24</b>	<b>80.88</b>	<b>71.36</b>	<b>26.98</b>	<b>39.82</b>	<b>27.65</b>	
Enc-dec, shared	Denoising	$P$	$M$	82.81	18.78	<b>80.63</b>	<b>70.73</b>	26.72	39.03	<b>27.46</b>	
Enc-dec, 6 layers	Denoising	$P$	$M/2$	80.88	18.97	77.59	68.42	26.38	38.40	26.95	
Language model	Denoising	$P$	$M$	74.70	17.93	61.14	55.02	25.09	35.28	25.86	
Prefix LM	Denoising	$P$	$M$	81.82	18.61	78.94	68.11	26.43	37.98	27.39	

Can we use a prefix LM model to bridge the gap?

? Other architectural choices: embeddings, activation functions, etc

---

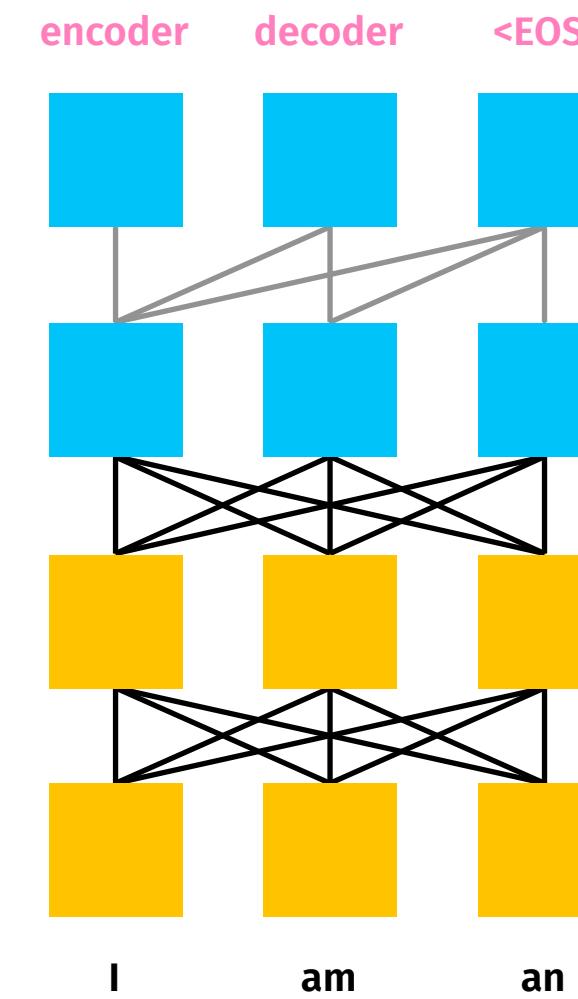
rotary, ALiBi

---

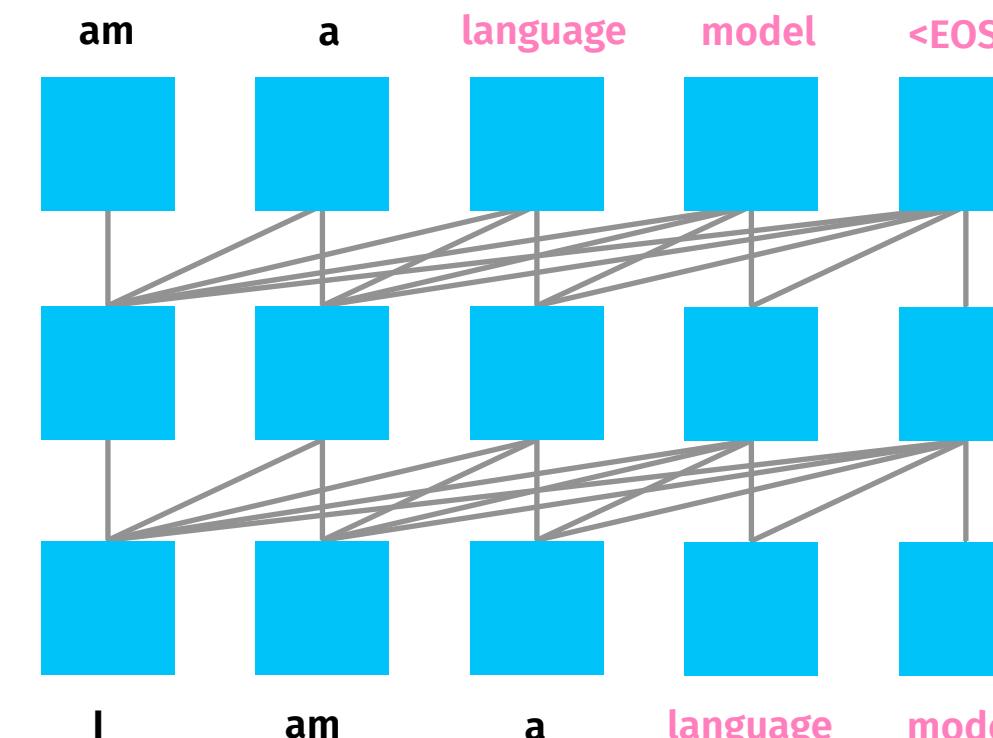
GeLU-GLU, squared ReLU

# Bridging the performance gap with prefix language modelling

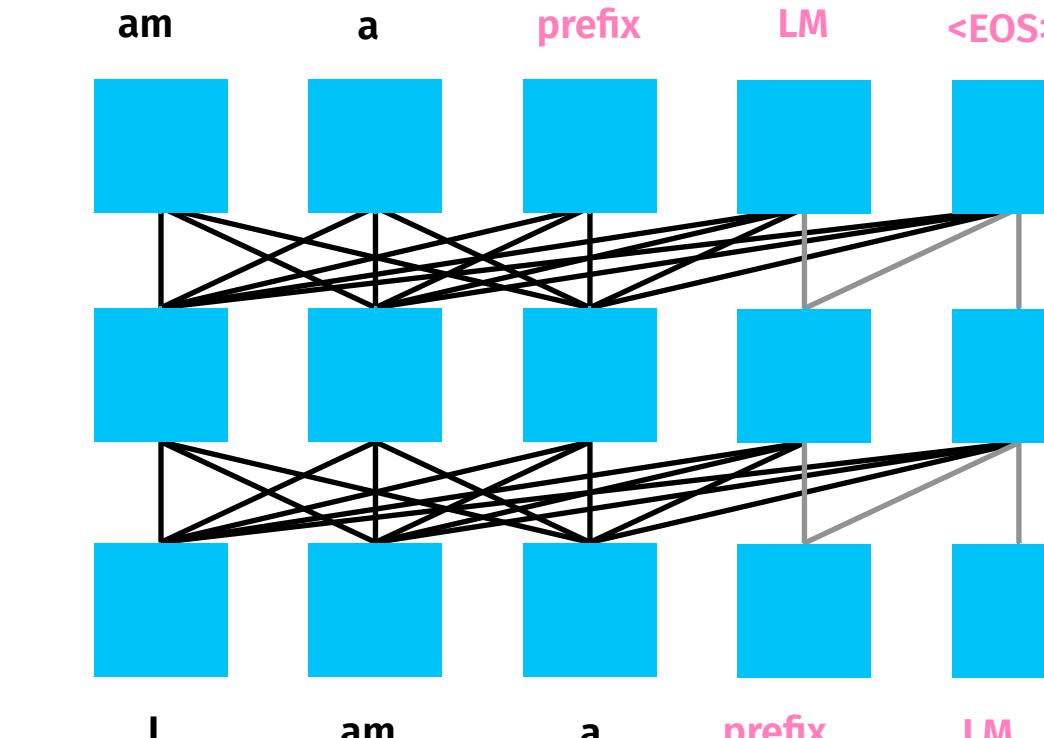
**encoder-decoder**  
e.g. T5



**autoregressive LM**  
e.g. GPT



**prefix LM**



encoder block

decoder block

— acausal attention

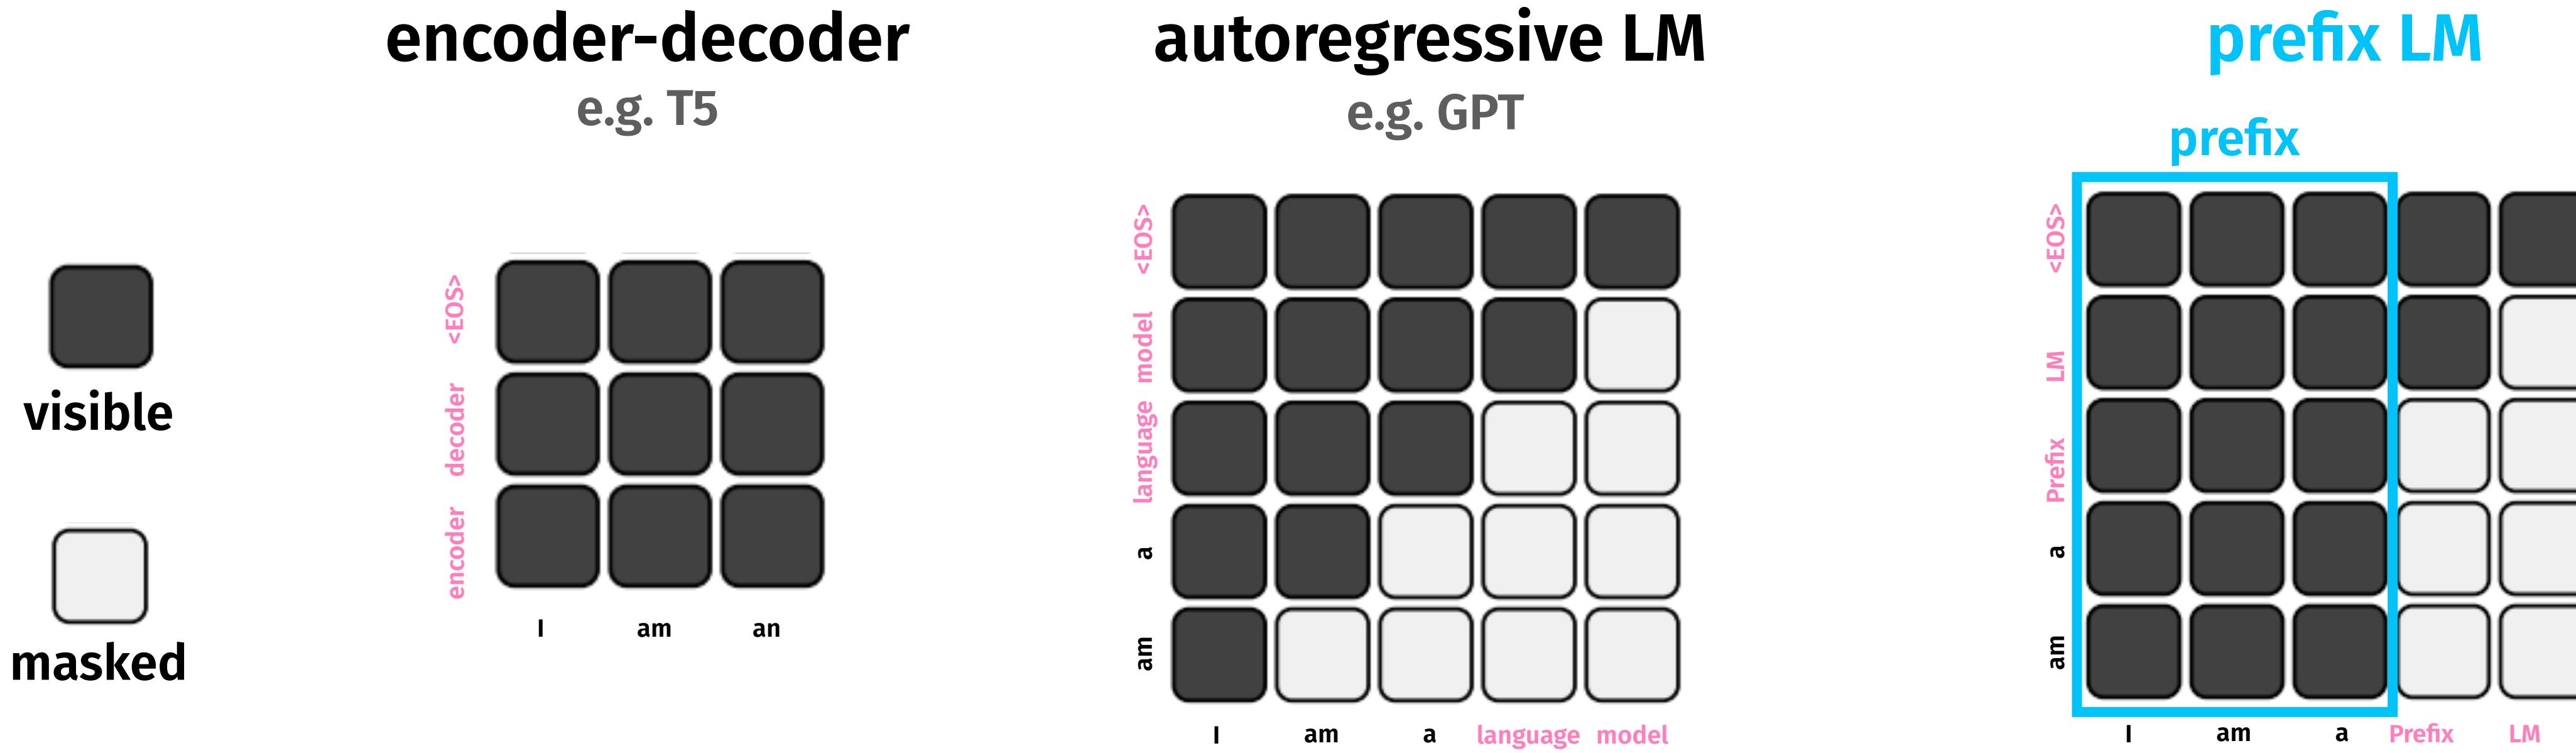
— causal attention

lorem ipsum  
lorem ipsum

**prompt text**  
**generated text**

⚠ Prefix LM: same architecture as autoregressive LM, but with a **different attention pattern**.

# Bridging the performance gap with prefix language modelling



💡 Intuition: tokens in the **prefix/prompt** don't have restricted view, thus better representation.

💡 As per T5, could **bridge encoder-decoder/LM gap**, but never demonstrated at scale nor for few-shot!

train with a randomly selected prefix during training, then prefix is prompt at inference time.

Megatron+DeepSpeed implementation ready, 1.3B results soon.

# Choosing a positional embedding: state-of-the-art



Better embeddings have been a hot topic: **rotary**, **ALiBi**, etc.

different metrics of importance: speed, stability, modeling loss, extrapolation.



## Rotary embeddings

clear performance advantage, very small cost in speed.

how it works: adds positional information to every layer, at the keys/queries.



## ALiBi: newest embedding, with extrapolation capabilities.

Extrapolation: pretrain on short sequences then evaluate on longer ones potentially opens the door to training with a smaller context size!

very simple and fast, performance on large models to be confirmed.

how it works: simple additive bias to attention scores

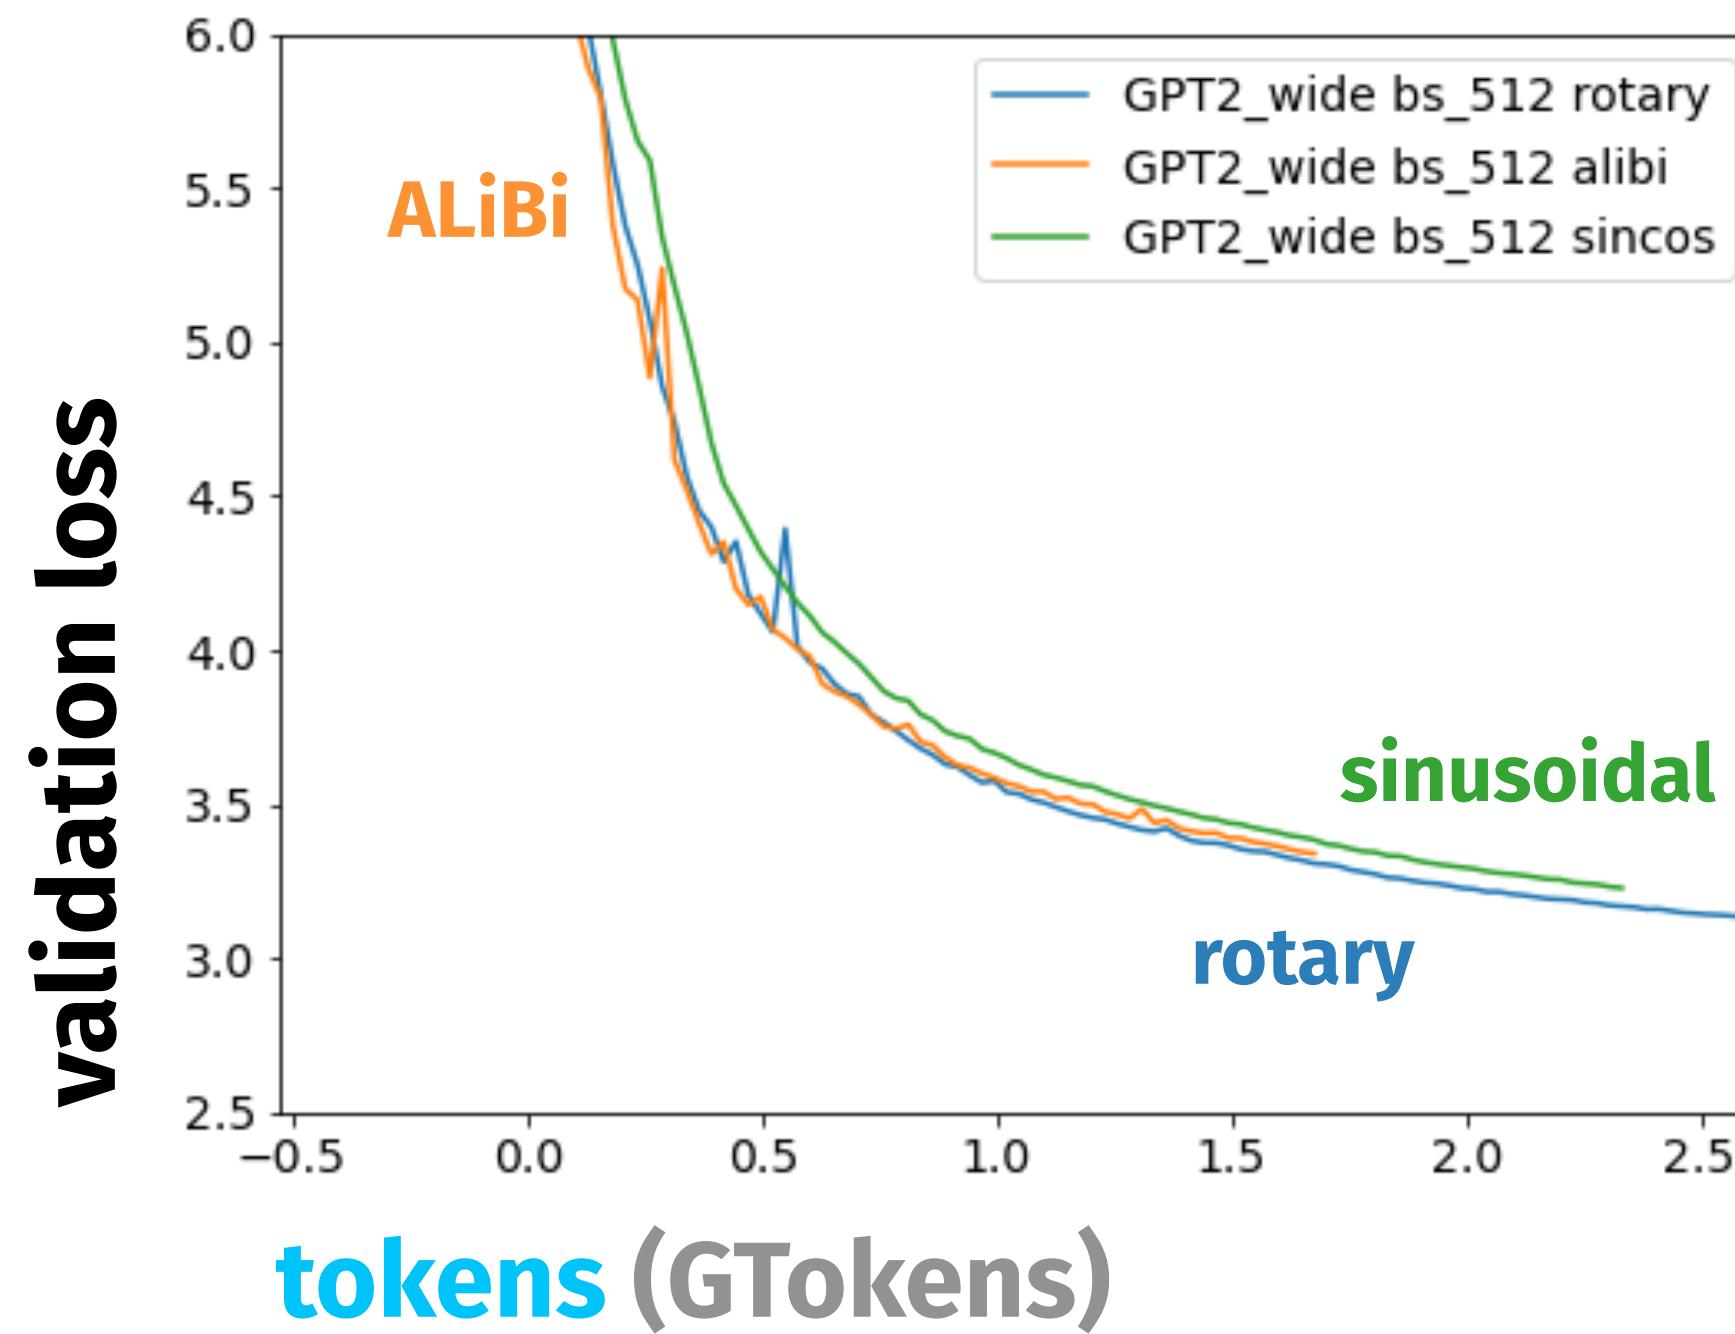
ALiBi, Press et al.

$$\begin{matrix} q_1 \cdot k_1 & & & & \\ & q_2 \cdot k_1 & q_2 \cdot k_2 & & \\ & q_3 \cdot k_1 & q_3 \cdot k_2 & q_3 \cdot k_3 & \\ & q_4 \cdot k_1 & q_4 \cdot k_2 & q_4 \cdot k_3 & q_4 \cdot k_4 \\ & q_5 \cdot k_1 & q_5 \cdot k_2 & q_5 \cdot k_3 & q_5 \cdot k_4 & q_5 \cdot k_5 \end{matrix} + \begin{matrix} 0 & & & & \\ -1 & 0 & & & \\ -2 & -1 & 0 & & \\ -3 & -2 & -1 & 0 & \\ -4 & -3 & -2 & -1 & 0 \end{matrix} \cdot m$$

# Choosing a positional embedding: first experiments

👍 **Rotary** and **ALiBi** consistently outperforms sinusoidal embeddings

why? they inject position information in each self-attention layer, not just in input embeddings;  
they use relative position information, so the model can't overfit certain locations.



**Limitations of evaluation so far:**

**medium model (350M) only, move to 1.3B**

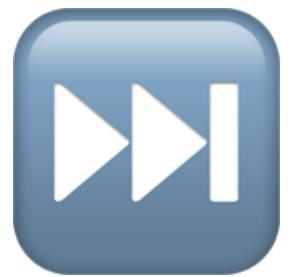
**LM loss only, should evaluate few-shot and more**

# Where we are and where are we going



done

**Implementation of different candidate architectures (mostly);**  
**Preprocessing multilingual training data;**  
**English-only baseline 1.3B run;**  
**English-only evaluation benchmark.**



next steps

**Debug training #1 (13B run) and understand few-shot performance;**  
**Evaluate English-only baseline on downstream tasks;**  
**Train and evaluate multilingual 1.3B baseline;**  
**Train and evaluate 1.3B ALiBi, rotary, and prefix LM.**

# 😊 Joining and contributing!

🌸 Join Big Science: <https://bigscience.huggingface.co/> and sign-up for modeling group.

🐙 GitHub: <https://github.com/bigscience-workshop/Megatron-DeepSpeed/issues>

 Weekly meetings: Wednesday 8am PT, 5pm CEST

## 😍 Contributors 😍



**Teven Le Scao**



**Sheng Shen**



**Thomas Wang**



**Ofir Press**



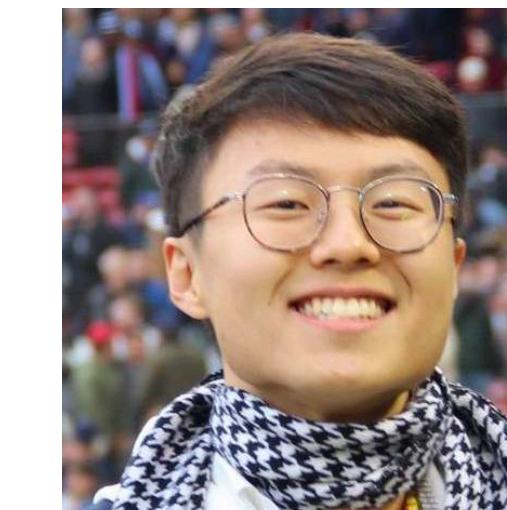
**Stella Biderman**



**M Saiful Bari (Maruf)**



**Lintang Sutawika**



**Jake Tae**



**Huu Nguyen**